

How far can a chain of surrogates get us?

Every expensive step in a structure-based pipeline handed to a learned surrogate. Other groups have done individual steps. We haven't found anyone who has run the whole chain, with the expensive physics kept only for the places where it changes the answer. It might not be good enough yet. We'd still like the number, as a baseline we'll come back to. **The testbed: can AI tell two near-identical drug targets apart?**

WHEN

{{DATES}}

Two days · 09:00 start

WHERE

{{VENUE}}

University of Auckland

THE GLOBAL EVENT

Runs 21-22 Oct, hubs on four continents.

Registration is open to anyone, come to that too.

WHY WE'RE DOING IT

Every step in the chain is improving quickly. At some point it crosses a threshold and becomes useful, and without a measured starting point we won't know when.

WHO WE NEED

Medicinal chemistry	Break our benchmark before we build on it	Structural biology	Judge predicted poses — we can't do this without you
Quantum chemistry	When is a cheap approximation good enough?	Statistics & ML	Bring your own small-data problem
Research software eng.	Make a GPU pipeline someone else can run	Any expensive science	Slow experiments, too many parameters — same problem

New to this?

Good. Within the first hour you'll own one configuration of the pipeline — one command, twenty minutes, and your run becomes a data point in the final analysis. Not a tutorial. The actual experiment. **Bring a laptop; that's all.**

WHAT YOU GET

Real hardware

Dual NVIDIA GB10 boxes and HGX H200 access — and a pipeline that already works when you walk in.

Your name on a result

Every participant's contribution lands in an artifact. Previous years produced papers crediting every team.

An honest answer

Six gates written before any data exists. A clean "no" is a successful outcome.