

Start small. Search hard. Scale only what survives.

We're testing whether a deliberately small experiment can predict what a large one would do. The testbed: **can AI tell two near-identical drug targets apart?** Most drugs fail on selectivity, not potency, and conventional screening can't separate a target from its close relatives — it flattens a 3D pocket into a fingerprint and discards exactly what makes them different. Two days to find out whether AI can put it back, with GPU quantum chemistry spent only where it changes the answer.

WHEN

{{DATES}}

Two days · 09:00 start

WHERE

{{VENUE}}

University of Auckland

TWO DAYS AHEAD OF THE WORLD

The global hackathon starts 21 Oct — we finish before the first hub anywhere opens.

WHY IT'S A REAL QUESTION

Everyone guesses most of their settings, because nobody can afford to explore them. Recent work says that's the mistake: small experiments fail because they're **undertuned**. Four configurations showed nothing. **256 gave a clean answer.**

WHO WE NEED

Medicinal chemistry	Break our benchmark before we build on it	Structural biology	Judge predicted poses — we can't do this without you
Quantum chemistry	When is a cheap approximation good enough?	Statistics & ML	Bring your own small-data problem
Research software eng.	Make a GPU pipeline someone else can run	Any expensive science	Slow experiments, too many parameters — same problem

New to this?

Good. Within the first hour you'll own one configuration of the pipeline — one command, twenty minutes, and your run becomes a data point in the final analysis. Not a tutorial. The actual experiment. **Bring a laptop; that's all.**

WHAT YOU GET

Real hardware Dual NVIDIA GB10 boxes and HGX H200 access — and a pipeline that already works when you walk in.	Your name on a result Every participant's contribution lands in an artifact. Previous years produced papers crediting every team.	An honest answer Six gates written before any data exists. A clean "no" is a successful outcome.
---	--	---